

18-213/18-613, Summer 2024
Attack Lab: Understanding Buffer Overflow Bugs
Assigned: Thursday, May 30
Due: Thursday, June 6, 11:59PM EDT
Last Possible Time to Turn in: Sunday, June 9, 11:59PM EDT

1 Introduction

This assignment involves generating a total of five attacks on two programs which have different security vulnerabilities. In this lab, you will:

- Learn different ways that attackers can exploit security vulnerabilities when programs do not safeguard themselves well enough against buffer overflows.
- Gain a better understanding of how to write programs that are more secure, as well as some of the features provided by compilers and operating systems to make programs less vulnerable.
- Gain a deeper understanding of the stack and parameter-passing mechanisms of x86-64 machine code.
- Gain a deeper understanding of how x86-64 instructions are encoded.
- Gain more experience with debugging tools such as GDB and OBJDUMP.

Note: In this lab, you will gain firsthand experience with methods used to exploit security weaknesses in operating systems and network servers. Our purpose is to help you learn about the runtime operation of programs and to understand the nature of these security weaknesses so that you can avoid them when you write system code. We do not condone the use of any other form of attack to gain unauthorized access to any system resources.

You will want to study Sections 3.10.3 and 3.10.4 of the book as reference material for this lab.

2 Logistics

As usual, this is an individual project. You will generate attacks for target programs that are custom generated for you.

2.1 Getting Files

You can obtain your files from the Autolab site

`https://ics.autolabproject.com`

After logging in to Autolab, select Attacklab -> Download handout. The Autolab server will build your files and return them to your browser in a tar file called `target k .tar`, where k is the unique number of your target programs.

Note: It takes about 15 seconds to build and download your target, so please be patient.

Save the `target k .tar` file in a (protected) Andrew directory in which you plan to do your work. Then login to a shark machine and give the command: `tar -xvf target k .tar`. This will extract a directory `target k` containing the files described below.

You should only download one set of files. If for some reason you download multiple targets, choose one target to work on and delete the rest.

Warning: If you expand your `target k .tar` on a PC, by using a utility such as Winzip, or letting your browser do the extraction, you'll risk resetting permission bits on the executable files.

The files in `target k` include:

`README.txt`: A file describing the contents of the directory

`ctarget`: An executable program vulnerable to *code-injection* attacks

`rtarget`: An executable program vulnerable to *return-oriented-programming* attacks

`cookie.txt`: An 8-digit hex code that you will use as a unique identifier in your attacks.

`farm.c`: The source code of your target's "gadget farm," which you will use in generating return-oriented programming attacks.

`hex2raw`: A utility to generate attack strings.

In the following instructions, we will assume that you have already copied the files to a protected local directory, and that you are executing the programs in that local directory.

2.2 Important Points

Here is a summary of some important rules regarding valid solutions for this lab. These points will not make much sense when you read this document for the first time. They are presented here as a central reference of rules once you get started.

- You must do the assignment on one of the class shark machines. There are ten machines available to students to use, a full list of which can be found on the course Web site at:

`http://www.cs.cmu.edu/~18213/labmachines.html`

- Your solutions may not use attacks to circumvent the validation code in the programs.
- You may only construct gadgets from file `rtarget` with addresses ranging between those for functions `start_farm` and `end_farm`.
- You can use any gadgets you discover in the farm, not just those for which we give the byte codes in Figure 3.
- You are allowed to use the standard tools for this assignment: `gdb` (or `lldb`), `objdump`, and any tool that you create (from scratch) on your own.
- You are **not** allowed to use tools or `gdb` plugins that are designed to assist in implementing buffer overflow and ROP attacks.
- It is **highly** recommended that you `gdb` to confirm that your exploit is placed on the stack with the correct byte ordering
- It is not uncommon for `gdb` to get *lost* while performing exploit code. You may be unable to step through your exploit in `gdb` while seeing all of the usual diagnostic information.

3 Target Programs

Both `CTARGET` and `RTARGET` read strings from standard input. They do so with the function `read_and_process_line` defined below:

```

1 /* read_and_process_line - Read a line from INFILE and process it.
2    This function has a buffer overflow vulnerability.  */
3 void read_and_process_line(FILE *infile) {
4     unsigned char buf[BUFFER_SIZE];
5     unsigned char *sp = buf;
6     int c;
7     while ((c = getc(infile)) != EOF && c != '\n') {
8         *sp++ = c;
9     }
10    process_line(buf, sp - buf);
11 }
```

This function reads a byte sequence from standard input, terminated by either a newline (`'\n'`), or end of file (EOF). Then it calls another function, `process_line`, passing it the bytes that were read. (It does *not* add a NUL terminator to the bytes it reads, so what it passes to `process_line` is not a “string.” Instead, it passes the number of bytes read as `process_line`’s second argument.)

In the code sample, you can see that `read_and_process_line` stores the byte sequence in a local variable `buf`, an array of `BUFFER_SIZE` bytes. (`BUFFER_SIZE` is a compile-time constant, specific to your version of `CTARGET` and `RTARGET`.) Notice that the `while` loop does not stop when `BUFFER_SIZE` bytes have been read. This is the same bug that’s found in the C library function `gets`: it keeps reading data until end of line or file, possibly overrunning the bounds of the storage allocated for the data.

If the input read by `read_and_process_line` is sufficiently short, nothing interesting will happen:

```
$ ./ctarget
Cookie: 0x599051eb
Type string:  Keep it short!
No exploit, read_and_process_line returned normally.
```

Typically an error occurs if you type a long string:

```
$unix ./ctarget
Cookie: 0x599051eb
Type string:  This is not a very interesting string, but it is quite long
Ouch!: You caused a segmentation fault!
Better luck next time
```

(Note that the value on the `Cookie:` line will differ from yours.)

of the cookie shown will differ from yours.) Program `RTARGET` will have the same behavior. As the error message indicates, overrunning the buffer typically causes the program state to be corrupted, leading to a memory access error. Your task is to be more clever with the strings you feed `CTARGET` and `RTARGET` so that they do more interesting things. These are called *exploit* strings.

Both `CTARGET` and `RTARGET` take several different command line arguments:

- h: Print list of possible command line arguments
- i FILE: Supply input from a file, rather than from standard input

Your exploit strings will typically contain byte values that do not correspond to the ASCII values for printing characters. The program `HEX2RAW` will enable you to generate these *raw* strings. See Appendix A for more information on how to use `HEX2RAW`.

Important points:

- Your exploit string must not contain byte value `0x0a` at any intermediate position, since this is the ASCII code for newline (`'\n'`). When `Gets` encounters this byte, it will assume you intended to terminate the string.
- `HEX2RAW` expects two-digit hex values separated by one or more white spaces. So if you want to create a byte with a hex value of 0, you need to write it as `00`. To create the word `0xdeadbeef` you should pass `"ef be ad de"` to `HEX2RAW` (note the reversal required for little-endian byte ordering).

When you have correctly solved one of the levels, your target program will automatically send a notification to Autolab. For example:

Phase	Program	Level	Method	Function	Points
1	CTARGET	1	CI	touch1	10
2	CTARGET	2	CI	touch2	25
3	CTARGET	3	CI	touch3	25
4	RTARGET	2	ROP	touch2	35
5	RTARGET	3	ROP	touch3	5

CI: Code injection

ROP: Return-oriented programming

Figure 1: Summary of attack lab phases

```
$unix ./hex2raw < ctarget.12.txt | ./ctarget
Cookie: 0x1a7dd803
Type string:Touch2!: You called touch2(0x1a7dd803)
Valid solution for level 2 with target ctarget
PASSED: Sent exploit string to server to be validated.
NICE JOB!
```

Unlike the Bomb Lab, there is no penalty for making mistakes in this lab. Feel free to fire away at CTARGET and RTARGET with any strings you like.¹

Figure 1 summarizes the five phases of the lab. As can be seen, the first three involve code-injection (CI) attacks on CTARGET, while the last two involve return-oriented-programming (ROP) attacks on RTARGET.

4 Part I: Code Injection Attacks

For the first three phases, your exploit strings will attack CTARGET. This program is set up in a way that the stack positions will be consistent from one run to the next and so that data on the stack can be treated as executable code. These features make the program vulnerable to attacks where the exploit strings contain the byte encodings of executable code.

4.1 Level 1

For Phase 1, you will not inject new code. Instead, your exploit string will redirect the program to execute an existing procedure.

Function `read_and_process_line` is called within CTARGET by a function `test` having the following C code:

```
1 void test(FILE *infile) {
2     read_and_process_line(infile);
```

¹Note that your target program will send notifications to the server for both successful and unsuccessful attacks. These data will enable us to track the overall progress of the class. Unsuccessful attacks will not have any effect on your score for this assignment.

```

3     notify_fail(0, "No exploit, read_and_process_line returned normally.");
4 }

```

When `read_and_process_line` executes its return statement (line 5 of `read_and_process_line`), the program ordinarily resumes execution within function `test` (at line 5 of this function). We want to change this behavior. Within the file `ctarget`, there is code for a function `touch1` having the following C representation:

```

1 TOUCH_FN touch1(void) {
2     vlevel = 1; /* Part of validation protocol */
3     printf("Touch1!: You called touch1()\n");
4     validate(1);
5 }

```

Your task is to get `CTARGET` to execute the code for `touch1` when `read_and_process_line` executes its return statement, rather than returning to `test`. Note that your exploit string may also corrupt parts of the stack not directly related to this stage, but this will not cause a problem, since `touch1` causes the program to exit directly.

Some Advice:

- All the information you need to devise your exploit string for this level can be determined by examining a disassembled version of `CTARGET`. Use `objdump -d` to get this disassembled version.
- The idea is to position a byte representation of the starting address for `touch1` so that the `ret` instruction at the end of the code for `read_and_process_line` will transfer control to `touch1`.
- Be careful about byte ordering. Remember, arrays (such as strings) are saved in index order, but values like integers are evaluated in little-endian.
- You might want to use GDB to step the program through the last few instructions of `read_and_process_line` to make sure it is doing the right thing.
- The placement of `buf` within the stack frame for `read_and_process_line` depends on the value of compile-time constant `BUFFER_SIZE`, as well the allocation strategy used by GCC. You will need to examine the disassembled code to determine its position.

4.2 Level 2

Phase 2 involves injecting a small amount of code as part of your exploit string.

Within the file `ctarget` there is code for a function `touch2` having the following C representation:

```

1 TOUCH_FN touch2(unsigned val) {
2     vlevel = 2; /* Part of validation protocol */
3     if (val == cookie) {
4         printf("Touch2!: You called touch2(0x%.8x)\n", val);

```

```

5         validate(2);
6     } else {
7         notify_fail(2, "Misfire: You called touch2(0x%.8x)", val);
8     }
9 }

```

Your task is to get CTARGET to execute the code for `touch2` rather than returning to `test`. In this case, however, you must make it appear to `touch2` as if you have passed your cookie as its argument.

Some Advice:

- You will want to position a byte representation of the address of your injected code in such a way that `ret` instruction at the end of the code for `read_and_process_line` will transfer control to it.
- Recall that the first argument to a function is passed in register `%rdi`.
- Your injected code should set the register to your cookie, and then use a `ret` instruction to transfer control to the first instruction in `touch2`.
- Do not attempt to use `jmp` or `call` instructions in your exploit code. The encodings of destination addresses for these instructions are difficult to formulate. Use `ret` instructions for all transfers of control, even when you are not returning from a call.
- See the discussion in Appendix B on how to use tools to generate the byte-level representations of instruction sequences.

4.3 Level 3

Phase 3 also involves a code injection attack, but passing a string as argument.

Within the file `ctarget` there is code for functions `hexmatch` and `touch3` having the following C representations:

```

1 /* Compare string to hex representation of unsigned value */
2 static int hexmatch(unsigned val, char *sval) {
3     char *endp;
4     unsigned long cval = strtoul(sval, &endp, 16);
5     return (cval == (unsigned long) val
6         && endp != sval
7         && *endp == '\0');
8 }
9
10 TOUCH_FN touch3(char *sval) {
11     vlevel = 3; /* Part of validation protocol */
12     if (hexmatch(cookie, sval)) {
13         report_touch3("Touch3!", sval);
14         validate(3);
15     } else {

```

```

16         report_touch3("Misfire", sval);
17         notify_fail(3, "touch3 called with the wrong cookie");
18     }
19 }

```

Your task is to get CTARGET to execute the code for `touch3` rather than returning to `test`. You must make it appear to `touch3` as if you have passed a string representation of your cookie as its argument.

Some Advice:

- You will need to include a string representation of your cookie in your exploit string. The string should consist of the eight hexadecimal digits (ordered from most to least significant) without a leading “0x.”
- Recall that a string is represented in C as a sequence of bytes followed by a byte with value 0. Type “`man ascii`” on any Linux machine to see the byte representations of the characters you need.
- Your injected code should set register `%rdi` to the address of this string.
- When functions `hexmatch` and `strncmp` are called, they push data onto the stack, overwriting portions of memory that held the buffer used by `read_and_process_line`. As a result, you will need to be careful where you place the string representation of your cookie.

5 Part II: Return-Oriented Programming

Performing code-injection attacks on program RTARGET is much more difficult than it is for CTARGET, because it uses two techniques to thwart such attacks:

- It uses randomization so that the stack positions differ from one run to another. This makes it impossible to determine where your injected code will be located.
- It marks the section of memory holding the stack as nonexecutable, so even if you could set the program counter to the start of your injected code, the program would fail with a segmentation fault.

Fortunately, clever people have devised strategies for getting useful things done in a program by executing existing code, rather than injecting new code. The most general form of this is referred to as *return-oriented programming* (ROP) [1, 2]. The strategy with ROP is to identify byte sequences within an existing program that consist of one or more instructions followed by the instruction `ret`. Such a segment is referred to as a *gadget*. Figure 2 illustrates how the stack can be set up to execute a sequence of n gadgets. In this figure, the stack contains a sequence of gadget addresses. Each gadget consists of a series of instruction bytes, with the final one being `0xc3`, encoding the `ret` instruction. When the program executes a `ret` instruction starting with this configuration, it will initiate a chain of gadget executions, with the `ret` instruction at the end of each gadget causing the program to jump to the beginning of the next.

A gadget can make use of code corresponding to assembly-language statements generated by the compiler, especially ones at the ends of functions. In practice, there may be some useful gadgets of this form, but not enough to implement many important operations. For example, it is highly unlikely that a compiled function




Figure 2: Setting up sequence of gadgets for execution. Byte value `0xc3` encodes the `ret` instruction.

would have `popq %rdi` as its last instruction before `ret`. Fortunately, with a byte-oriented instruction set, such as x86-64, a gadget can often be found by extracting patterns from other parts of the instruction byte sequence.

For example, one version of `rtarget` contains code generated for the following C function:

```
void setval_210(unsigned *p)
{
    *p = 3347663060U;
}
```

The chances of this function being useful for attacking a system seem pretty slim. But, the disassembled machine code for this function shows an interesting byte sequence:

```
0000000000400f15 <setval_210>:
    400f15:    c7 07 d4 48 89 c7    movl    $0xc78948d4, (%rdi)
    400f1b:    c3                  retq
```

The byte sequence `48 89 c7` encodes the instruction `movq %rax, %rdi`. (See Figure 3A for the encodings of useful `movq` instructions.) This sequence is followed by byte value `c3`, which encodes the `ret` instruction. The function starts at address `0x400f15`, and the sequence starts on the fourth byte of the function. Thus, this code contains a gadget, having a starting address of `0x400f18`, that will copy the 64-bit value in register `%rax` to register `%rdi`.

Your code for `RTARGET` contains a number of functions similar to the `setval_210` function shown above in a region we refer to as the *gadget farm*. Your job will be to identify useful gadgets in the gadget farm and use these to perform attacks similar to those you did in Phases 2 and 3.

Important: The gadget farm is demarcated by functions `start_farm` and `end_farm` in your copy of `rtarget`. Do not attempt to construct gadgets from other portions of the program code.

A. Encodings of movq instructions

movq S, D

Source S	Destination D							
	%rax	%rcx	%rdx	%rbx	%rsp	%rbp	%rsi	%rdi
%rax	48 89 c0	48 89 c1	48 89 c2	48 89 c3	48 89 c4	48 89 c5	48 89 c6	48 89 c7
%rcx	48 89 c8	48 89 c9	48 89 ca	48 89 cb	48 89 cc	48 89 cd	48 89 ce	48 89 cf
%rdx	48 89 d0	48 89 d1	48 89 d2	48 89 d3	48 89 d4	48 89 d5	48 89 d6	48 89 d7
%rbx	48 89 d8	48 89 d9	48 89 da	48 89 db	48 89 dc	48 89 dd	48 89 de	48 89 df
%rsp	48 89 e0	48 89 e1	48 89 e2	48 89 e3	48 89 e4	48 89 e5	48 89 e6	48 89 e7
%rbp	48 89 e8	48 89 e9	48 89 ea	48 89 eb	48 89 ec	48 89 ed	48 89 ee	48 89 ef
%rsi	48 89 f0	48 89 f1	48 89 f2	48 89 f3	48 89 f4	48 89 f5	48 89 f6	48 89 f7
%rdi	48 89 f8	48 89 f9	48 89 fa	48 89 fb	48 89 fc	48 89 fd	48 89 fe	48 89 ff

B. Encodings of popq instructions

Operation	Register R							
	%rax	%rcx	%rdx	%rbx	%rsp	%rbp	%rsi	%rdi
popq R	58	59	5a	5b	5c	5d	5e	5f

C. Encodings of movl instructions

movl S, D

Source S	Destination D							
	%eax	%ecx	%edx	%ebx	%esp	%ebp	%esi	%edi
%eax	89 c0	89 c1	89 c2	89 c3	89 c4	89 c5	89 c6	89 c7
%ecx	89 c8	89 c9	89 ca	89 cb	89 cc	89 cd	89 ce	89 cf
%edx	89 d0	89 d1	89 d2	89 d3	89 d4	89 d5	89 d6	89 d7
%ebx	89 d8	89 d9	89 da	89 db	89 dc	89 dd	89 de	89 df
%esp	89 e0	89 e1	89 e2	89 e3	89 e4	89 e5	89 e6	89 e7
%ebp	89 e8	89 e9	89 ea	89 eb	89 ec	89 ed	89 ee	89 ef
%esi	89 f0	89 f1	89 f2	89 f3	89 f4	89 f5	89 f6	89 f7
%edi	89 f8	89 f9	89 fa	89 fb	89 fc	89 fd	89 fe	89 ff

D. Encodings of 2-byte functional nop instructions

Operation		Register R			
		%al	%cl	%dl	%bl
andb	R, R	20 c0	20 c9	20 d2	20 db
orb	R, R	08 c0	08 c9	08 d2	08 db
cmpb	R, R	38 c0	38 c9	38 d2	38 db
testb	R, R	84 c0	84 c9	84 d2	84 db

Figure 3: Byte encodings of instructions. All values are shown in hexadecimal.

5.1 Level 2

For Phase 4, you will repeat the attack of Phase 2, but do so on program `RTARGET` using gadgets from your gadget farm. You can construct your solution using gadgets consisting of the following instruction types, and using only the first eight x86-64 registers (`%rax-%rdi`).

`movq` : The codes for these are shown in Figure 3A.

`popq` : The codes for these are shown in Figure 3B.

`ret` : This instruction is encoded by the single byte `0xc3`.

`nop` : This instruction (pronounced “no op,” which is short for “no operation”) is encoded by the single byte `0x90`. Its only effect is to cause the program counter to be incremented by 1.

Some Advice:

- All the gadgets you need can be found in the region of the code for `rtarget` demarcated by the functions `start_farm` and `mid_farm`.
- You can do this attack with just two gadgets.
- When a gadget uses a `popq` instruction, it will pop data from the stack. As a result, your exploit string will contain a combination of gadget addresses and data.

5.2 Level 3

Before you take on the Phase 5, pause to consider what you have accomplished so far. In Phases 2 and 3, you caused a program to execute machine code of your own design. If `CTARGET` had been a network server, you could have injected your own code into a distant machine. In Phase 4, you circumvented two of the main devices modern systems use to thwart buffer overflow attacks. Although you did not inject your own code, you were able to inject a type of program that operates by stitching together sequences of existing code. You have also gotten 95/100 points for the lab. That’s a good score. If you have other pressing obligations consider stopping right now.

Phase 5 requires you to do an ROP attack on `RTARGET` to invoke function `touch3` with a pointer to a string representation of your cookie. That may not seem significantly more difficult than using an ROP attack to invoke `touch2`, except that we have made it so. Moreover, Phase 5 counts for only 5 points, which is not a true measure of the effort it will require. Think of it as more an extra credit problem for those who want to go beyond the normal expectations for the course.

To solve Phase 5, you can use gadgets in the region of the code in `rtarget` demarcated by functions `start_farm` and `end_farm`. In addition to the gadgets used in Phase 4, this expanded farm includes the encodings of different `movl` instructions, as shown in Figure 3C. The byte sequences in this part of the farm also contain 2-byte instructions that serve as *functional nops*, i.e., they do not change any register or memory values. These include instructions, shown in Figure 3D, such as `andb %al, %al`, that operate on the low-order bytes of some of the registers but do not change their values.

Some Advice:

- You'll want to review the effect a `movl` instruction has on the upper 4 bytes of a register, as is described on page 183 of the text.
- The official solution requires a sequence of eight gadgets. Depending on the contents of your farm (each target has a different one), you may be able find a shorter one.
- Remember: Your exploit string must not contain the newline character (byte value 0x0a) at any intermediate position

6 Logistical Notes

Whenever you correctly solve a phase, your target program will send a message to Autolab. The server will test your exploit string to make sure it really works, (From Autolab, follow `Attacklab -> View scoreboard`). This Web page is updated every minute or so to show the progress for each target. You should be sure to check this page after your submission to make sure your string has been validated. (If you really solved the phase, your string *should* be valid. But, an attack that works on your copy of one of the targets may be rejected by the version on the server, if it does not correctly implement the specified functionality.)

Each phase is graded individually. You do not need to do them in the specified order, but you will get credit only for the phases for which the server receives a valid message.

Good luck and have fun!

A Using HEX2RAW

HEX2RAW takes as input a *hex-formatted* string. In this format, each byte value is represented by two hex digits. For example, the string "012345" could be entered in hex format as "30 31 32 33 34 35 00." (Recall that the ASCII code for decimal digit x is $0 \times 3x$, and that the end of a string is indicated by a null byte.)

The hex characters you pass to HEX2RAW should be separated by whitespace (blanks or newlines). We recommend separating different parts of your exploit string with newlines while you're working on it. HEX2RAW supports C-style block comments, so you can mark off sections of your exploit string. For example:

```
48 c7 c1 f0 11 40 00 /* mov    $0x4011f0,%rcx */
```

Be sure to leave space around both the starting and ending comment strings ("`/*`", "`*/`"), so that the comments will be properly ignored.

If you generate a hex-formatted exploit string in the file `exploit.txt`, you can apply the raw string to CTARGET or RTARGET in several different ways:

1. You can set up a series of pipes to pass the string through HEX2RAW.

```
$unix cat exploit.txt | ./hex2raw | ./ctarget
```

2. You can store the raw string in a file and use I/O redirection:

```
$unix ./hex2raw exploit.txt exploit-raw.txt
$unix ./ctarget < exploit-raw.txt
```

This approach can also be used when running from within GDB:

```
$unix gdb ctarget
(gdb) run < exploit-raw.txt
```

3. You can store the raw string in a file and provide the file name as a command-line argument:

```
$unix ./hex2raw exploit.txt exploit-raw.txt
$unix ./ctarget -i exploit-raw.txt
```

This approach also can be used when running from within GDB.

B Generating Byte Codes

Using GCC as an assembler and OBJDUMP as a disassembler makes it convenient to generate the byte codes for instruction sequences. For example, suppose you write a file `example.s` containing the following assembly code:

```
# Example of hand-generated assembly code
pushq    $0xabcdef          # Push value onto stack
addq     $17,%rax            # Add 17 to %rax
movl     %eax,%edx           # Copy lower 32 bits to %edx
```

The code can contain a mixture of instructions and data. Anything to the right of a '#' character is a comment.

You can now assemble and disassemble this file:

```
$unix gcc -c example.s
$unix objdump -d example.o > example.d
```

The generated file `example.d` contains the following:

```
example.o:      file format elf64-x86-64
```

Disassembly of section `.text`:

```

000000000000000000 <.text>:
  0: 68 ef cd ab 00      push    $0xabcdef
  5: 48 83 c0 11      add     $0x11,%rax
  9: 89 c2              mov     %eax,%edx

```

The lines at the bottom show the machine code generated from the assembly language instructions. Each line has a hexadecimal number on the left indicating the instruction's starting address (starting with 0), while the hex digits after the ':' character indicate the byte codes for the instruction. Thus, we can see that the instruction `push $0xABCDEF` has hex-formatted byte code `68 ef cd ab 00`.

From this file, you can get the byte sequence for the code:

```
68 ef cd ab 00 48 83 c0 11 89 c2
```

This string can then be passed through `HEX2RAW` to generate an input string for the target programs.. Alternatively, you can edit `example.d` to omit extraneous values and to contain C-style comments for readability, yielding:

```

68 ef cd ab 00    /* pushq    $0xabcdef */
48 83 c0 11       /* add     $0x11,%rax */
89 c2             /* mov     %eax,%edx */

```

This is also a valid input you can pass through `HEX2RAW` before sending to one of the target programs.

C Extra Reading

You can find more on vulnerability exploits with this light reading:

<https://blog.skullsecurity.org/2013/ropasaurusrex-a-primer-on-return-oriented-programming>

References

- [1] R. Roemer, E. Buchanan, H. Shacham, and S. Savage. Return-oriented programming: Systems, languages, and applications. *ACM Transactions on Information System Security*, 15(1):2:1–2:34, March 2012.
- [2] E. J. Schwartz, T. Avgerinos, and D. Brumley. Q: Exploit hardening made easy. In *USENIX Security Symposium*, 2011.